

---

# **Survey of DNN Hardware**

# Hardware are targeting deep learning

---

- CPU
  - Intel Knights Landing
  - Intel Knight Mill
- GPU
  - PASCAL
  - VOLTA
- System for deep learning
  - Nvidia DGX-1 (2016)
- Cloud Systems for Deep Learning
- SOCs for Deep Learning Inference
- FPGAs for Deep Learning

# CPUs Are Targeting Deep Learning

## Intel Knights Landing (2016)



- 7 TFLOPS FP32
- 16GB MCDRAM– 400 GB/s
- 245W TDP
- 29 GFLOPS/W (FP32)
- 14nm process

## **Knights Mill:** next gen Xeon Phi “optimized for deep learning”

Intel announced the addition of **new vector instructions** for deep learning (AVX512-4VNNIW and AVX512-4FMAPS), October 2016

# CPUs Are Targeting Deep Learning

---

- Knights Mill
  - Intel's codename for a Xeon Phi product specialized in **deep learning**,
  - Initially released in December 2017
  - Knights Mill includes optimizations for better utilization of AVX-512 instructions and enables 4-way hyperthreading.
  - **Single-precision** and **variable-precision floating-point** performance increased.



# GPUs Are Targeting Deep Learning

## Nvidia PASCAL GP100 (2016)




- **10/20 TFLOPS FP32/FP16**
- 16GB HBM – 750 GB/s
- 300W TDP
- 33/67 GFLOPS/W (FP32/FP16)
- 16nm process
- 160GB/s NV Link

FP16 support to perform two FP16 operations on a single precision core for faster deep learning computation

# GPUs Are Targeting Deep Learning

## Nvidia VOLTA GV100 (2017)



- 15 TFLOPS FP32
- 16GB HBM2 – 900 GB/s
- 300W TDP
- 50 GFLOPS/W (FP32)
- 12nm process
- 300GB/s NV Link2

Tensor Core....

Source: Nvidia

# GV100 – “Tensor Core”

---

$$\mathbf{D} = \begin{pmatrix} A_{0,0} & A_{0,1} & A_{0,2} & A_{0,3} \\ A_{1,0} & A_{1,1} & A_{1,2} & A_{1,3} \\ A_{2,0} & A_{2,1} & A_{2,2} & A_{2,3} \\ A_{3,0} & A_{3,1} & A_{3,2} & A_{3,3} \end{pmatrix} \begin{pmatrix} B_{0,0} & B_{0,1} & B_{0,2} & B_{0,3} \\ B_{1,0} & B_{1,1} & B_{1,2} & B_{1,3} \\ B_{2,0} & B_{2,1} & B_{2,2} & B_{2,3} \\ B_{3,0} & B_{3,1} & B_{3,2} & B_{3,3} \end{pmatrix} + \begin{pmatrix} C_{0,0} & C_{0,1} & C_{0,2} & C_{0,3} \\ C_{1,0} & C_{1,1} & C_{1,2} & C_{1,3} \\ C_{2,0} & C_{2,1} & C_{2,2} & C_{2,3} \\ C_{3,0} & C_{3,1} & C_{3,2} & C_{3,3} \end{pmatrix}$$

FP16 or FP32                      FP16                      FP16                      FP16 or FP32

Efficient Execution of 4x4 FP16 Multiplication and Addition

Tensor Core....

- 120 TFLOPS (FP16)
- 400 GFLOPS/W (FP16)

# Systems for Deep Learning

---

## Nvidia DGX-2 (2018)



- 2 peta FLOPS
- 16× Tesla V100, Dual Xeon
- 512GB GPU Memory
- 12 NVIDIA NVSwitch
- Optimized DL Software
- 7 TB SSD Storage
- Dual 10GbE, 8X 100Gb
- 10000W



# Cloud Systems for Deep Learning

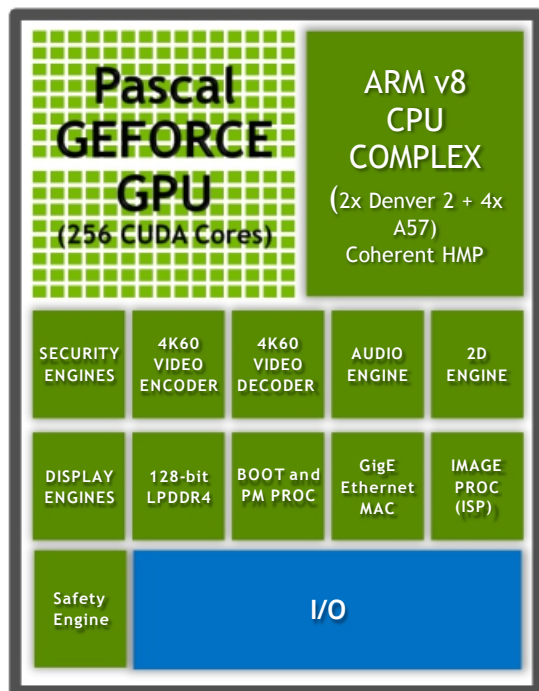
## Facebook's Deep Learning Machine



- Open Rack Compliant
- Powered by 8 Tesla M40 GPUs
- 2x Faster Training for Faster Deployment
- 2x Larger Networks for Higher Accuracy

# SOCs for Deep Learning Inference

## Nvidia Tegra - Parker




- GPU: 1.5 TeraFLOPS FP16
- 4GB LPDDR4 @ 25.6 GB/s
- 15 W TDP  
(1W idle, <10W typical)
- 100 GFLOPS/W (FP16)
- 16nm process

**Xavier:** next gen Tegra to be an “AI supercomputer”

# Mobile SOCs for Deep Learning

- Samsung Exynos (ARM Mali)
  - Exynos 8 Octa 8890
  - GPU: 0.26 TFLOPS
  - LPDDR4 @ 28.7 GB/s
  - 14nm process
- Source: Wikipedia
- Newer version
  - Exynos 9 Octa 8895 (S9/S9+)
  - Exynos 9 Octa 9820 (S10/S10+)



# FPGAs for Deep Learning

---



## Intel/Altera Stratix 10

- 10 TFLOPS FP32
- HBM2 integrated
- Up to 1 GHz
- 14nm process
- 80 GFLOPS/W



## Xilinx Virtex UltraSCALE+

- DSP: up to 21.2 TMACS
- DSP: up to 890 MHz
- Up to 500Mb On-Chip Memory
- 16nm process

---

# **Kernel Computation**

# Convolution (CONV) Layer

- Convert to matrix mult. using the **Toeplitz Matrix**

Convolution:

$$\begin{array}{|c|c|} \hline 1 & 2 \\ \hline 3 & 4 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 2 & 3 \\ \hline 4 & 5 & 6 \\ \hline 7 & 8 & 9 \\ \hline \end{array} = \begin{array}{|c|c|} \hline 1 & 2 \\ \hline 3 & 4 \\ \hline \end{array}$$



Matrix Mult:

**Toeplitz Matrix  
(w/ redundant data)**

$$\begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline \end{array} \times \begin{array}{|c|c|c|c|} \hline 1 & 2 & 4 & 5 \\ \hline 2 & 3 & 5 & 6 \\ \hline 4 & 5 & 7 & 8 \\ \hline 5 & 6 & 8 & 9 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 1 & 2 & 3 & 4 \\ \hline \end{array}$$

# Convolution (CONV) Layer

- Convert to matrix mult. using the **Toeplitz Matrix**

Convolution:

Filter		Input Fmap		Output Fmap																	
<table><tr><td>1</td><td>2</td></tr><tr><td>3</td><td>4</td></tr></table>	1	2	3	4	*	<table><tr><td>1</td><td>2</td><td>3</td></tr><tr><td>4</td><td>5</td><td>6</td></tr><tr><td>7</td><td>8</td><td>9</td></tr></table>	1	2	3	4	5	6	7	8	9	=	<table><tr><td>1</td><td>2</td></tr><tr><td>3</td><td>4</td></tr></table>	1	2	3	4
1	2																				
3	4																				
1	2	3																			
4	5	6																			
7	8	9																			
1	2																				
3	4																				



Matrix Mult:

		Toeplitz Matrix (w/ redundant data)																											
	<table><tr><td>1</td><td>2</td><td>3</td><td>4</td></tr></table>	1	2	3	4	×	<table><tr><td>1</td><td>2</td><td>4</td><td>5</td></tr><tr><td>2</td><td>3</td><td>5</td><td>6</td></tr><tr><td>4</td><td>5</td><td>7</td><td>8</td></tr><tr><td>5</td><td>6</td><td>8</td><td>9</td></tr></table>	1	2	4	5	2	3	5	6	4	5	7	8	5	6	8	9	=	<table><tr><td>1</td><td>2</td><td>3</td><td>4</td></tr></table>	1	2	3	4
1	2	3	4																										
1	2	4	5																										
2	3	5	6																										
4	5	7	8																										
5	6	8	9																										
1	2	3	4																										

Data is repeated

---

# **Computational Transforms**

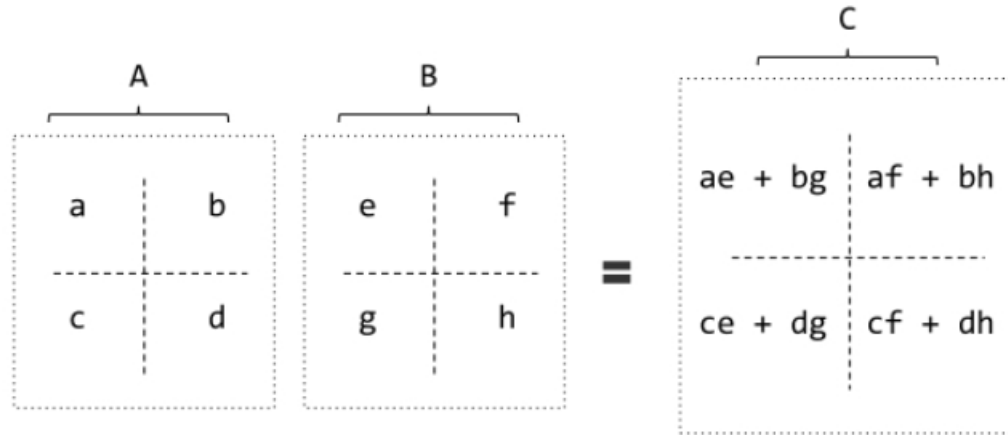


# Computation Transformations

---

- **Goal: Bitwise same result, but reduce number of operations**
- **Focuses mostly on compute**

# Strassen



8 multiplications + 4 additions

$$\begin{aligned} P1 &= a(f - h) \\ P2 &= (a + b)h \\ P3 &= (c + d)e \\ P4 &= d(g - e) \end{aligned}$$

$$\begin{aligned} P5 &= (a + d)(e + h) \\ P6 &= (b - d)(g + h) \\ P7 &= (a - c)(e + f) \end{aligned}$$

$$AB = \begin{bmatrix} P5 + P4 - P2 + P6 & P1 + P2 \\ P3 + P4 & P1 + P5 - P3 - P7 \end{bmatrix}$$


7 multiplications + 18 additions

[Cong et al., ICANN, 2014]

# Strassen

---

- Reduce the complexity of matrix multiplication from  $\Theta(N^3)$  to  $\Theta(N^{2.807})$  by reducing multiplication



Comes at the price of reduced numerical stability  
and requires significantly more memory

# Winograd 1D – F(2,3)

---

- Targeting convolutions instead of matrix multiply
- Notation: F(size of output, filter size)

$$F(2, 3) = \begin{matrix} & \text{input} & & \text{filter} \\ \begin{bmatrix} d_0 & d_1 & d_2 \\ d_1 & d_2 & d_3 \end{bmatrix} & \begin{bmatrix} g_0 \\ g_1 \\ g_2 \end{bmatrix} & = & \begin{bmatrix} d_0 g_0 + d_0 g_1 + d_0 g_2 \\ d_1 g_0 + d_1 g_1 + d_1 g_2 \end{bmatrix} \end{matrix}$$

6 multiplications + 4 additions

# Winograd 1D – F(2,3)

- Targeting convolutions instead of matrix multiply
- Notation: F(size of output, filter size)

$$F(2, 3) = \begin{matrix} & \text{input} & & \text{filter} \\ \begin{bmatrix} d_0 & d_1 & d_2 \\ d_1 & d_2 & d_3 \end{bmatrix} & \begin{bmatrix} g_0 \\ g_1 \\ g_2 \end{bmatrix} & = & \begin{bmatrix} m_1 + m_2 + m_3 \\ m_2 - m_3 - m_4 \end{bmatrix} \end{matrix}$$

$$\begin{aligned} m_1 &= (d_0 - d_2)g_0 & m_2 &= (d_1 + d_2) \frac{g_0 + g_1 + g_2}{2} \\ m_4 &= (d_1 - d_3)g_2 & m_3 &= (d_2 - d_1) \frac{g_0 - g_1 + g_2}{2} \end{aligned}$$

4 multiplications + 12 additions + 2 shifts

# Winograd 2D - F(2x2, 3x3)

- 1D Winograd is nested to make 2D Winograd

Filter		Input Fmap		Output Fmap																													
<table><tr><td>g<sub>00</sub></td><td>g<sub>01</sub></td><td>g<sub>02</sub></td></tr><tr><td>g<sub>10</sub></td><td>g<sub>11</sub></td><td>g<sub>12</sub></td></tr><tr><td>g<sub>20</sub></td><td>g<sub>21</sub></td><td>g<sub>22</sub></td></tr></table>	g <sub>00</sub>	g <sub>01</sub>	g <sub>02</sub>	g <sub>10</sub>	g <sub>11</sub>	g <sub>12</sub>	g <sub>20</sub>	g <sub>21</sub>	g <sub>22</sub>	*	<table><tr><td>d<sub>00</sub></td><td>d<sub>01</sub></td><td>d<sub>02</sub></td><td>d<sub>03</sub></td></tr><tr><td>d<sub>10</sub></td><td>d<sub>11</sub></td><td>d<sub>12</sub></td><td>d<sub>13</sub></td></tr><tr><td>d<sub>20</sub></td><td>d<sub>21</sub></td><td>d<sub>22</sub></td><td>d<sub>23</sub></td></tr><tr><td>d<sub>30</sub></td><td>d<sub>31</sub></td><td>d<sub>32</sub></td><td>d<sub>33</sub></td></tr></table>	d <sub>00</sub>	d <sub>01</sub>	d <sub>02</sub>	d <sub>03</sub>	d <sub>10</sub>	d <sub>11</sub>	d <sub>12</sub>	d <sub>13</sub>	d <sub>20</sub>	d <sub>21</sub>	d <sub>22</sub>	d <sub>23</sub>	d <sub>30</sub>	d <sub>31</sub>	d <sub>32</sub>	d <sub>33</sub>	=	<table><tr><td>y<sub>00</sub></td><td>y<sub>01</sub></td></tr><tr><td>y<sub>10</sub></td><td>y<sub>11</sub></td></tr></table>	y <sub>00</sub>	y <sub>01</sub>	y <sub>10</sub>	y <sub>11</sub>
g <sub>00</sub>	g <sub>01</sub>	g <sub>02</sub>																															
g <sub>10</sub>	g <sub>11</sub>	g <sub>12</sub>																															
g <sub>20</sub>	g <sub>21</sub>	g <sub>22</sub>																															
d <sub>00</sub>	d <sub>01</sub>	d <sub>02</sub>	d <sub>03</sub>																														
d <sub>10</sub>	d <sub>11</sub>	d <sub>12</sub>	d <sub>13</sub>																														
d <sub>20</sub>	d <sub>21</sub>	d <sub>22</sub>	d <sub>23</sub>																														
d <sub>30</sub>	d <sub>31</sub>	d <sub>32</sub>	d <sub>33</sub>																														
y <sub>00</sub>	y <sub>01</sub>																																
y <sub>10</sub>	y <sub>11</sub>																																

**Original:** 36 multiplications


**Winograd:** 16 multiplications → 2.25 times reduction

# Winograd Summary

---

- **Winograd is an optimized computation for convolutions**
- **It can significantly reduce multiplies**
  - **For example, for 3x3 filter by 2.25X**
- **But, each filter size is a different computation.**

# FFT Flow





# FFT Overview


---

- Convert **filter** and **input** to **frequency** domain to make **convolution** a simple **multiply** then convert back to **time** domain.
- Convert direct convolution  $O(N_o^2 N_f^2)$  computation to  $O(N_o^2 \log_2 N_o)$
- So note that computational **benefit** of FFT decreases with **decreasing size** of filter

[Mathieu et al., ArXiv 2013, Vasilache et al., ArXiv 2014]

# cuDNN: Speed up with Transformations

60x Faster Training in 3 Years



AlexNet training throughput on:

CPU: 1x E5-2680v3 12 Core 2.5GHz. 128GB System Memory, Ubuntu 14.04

M40 bar: 8x M40 GPUs in a node, P100: 8x P100 NVLink-enabled

Source: Nvidia

---

# Backup Slides